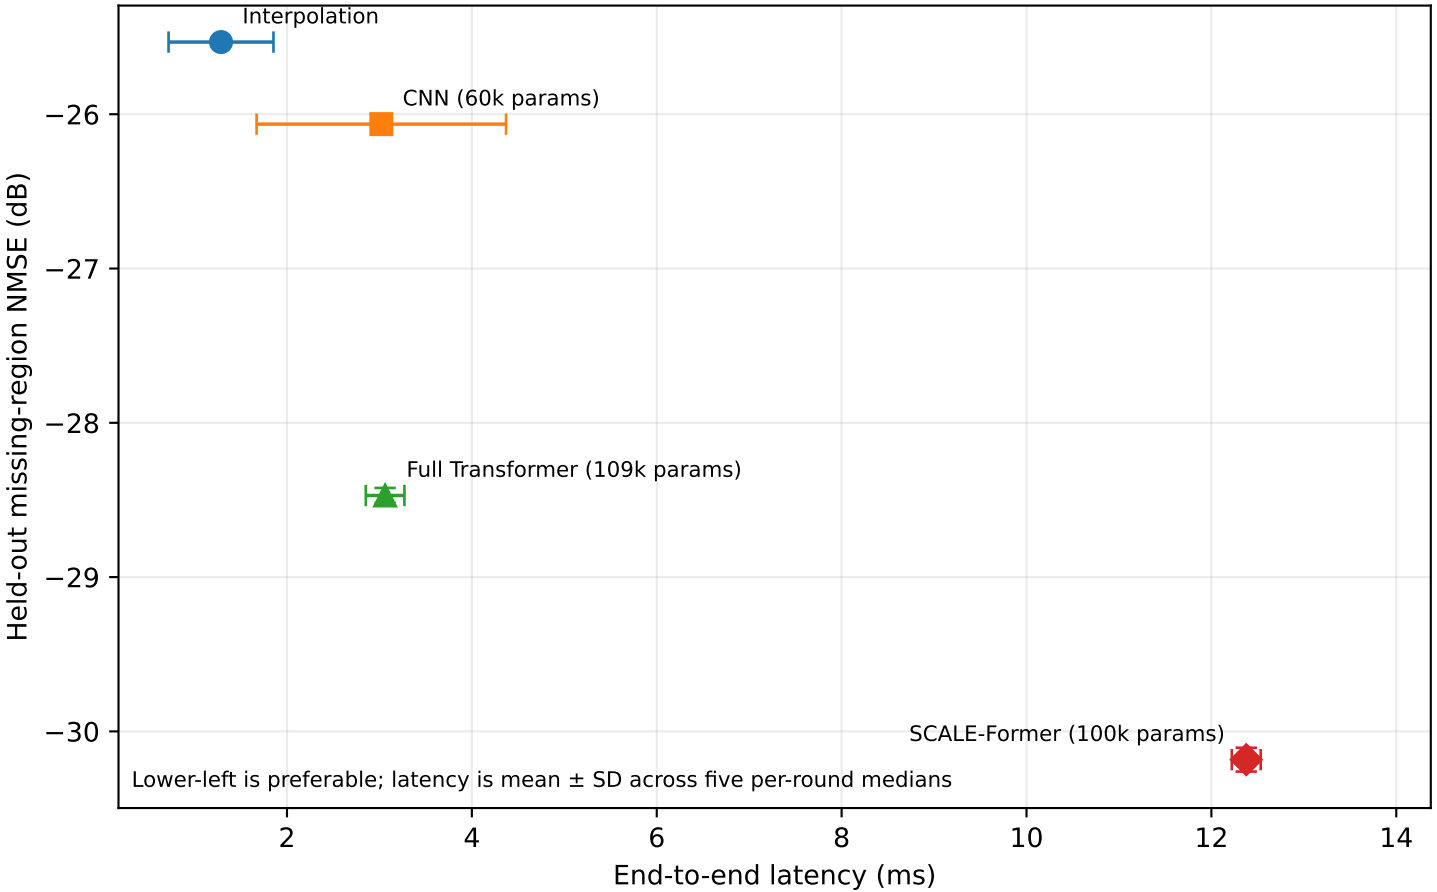
Held-out accuracy versus inference latency
Interpolation
CNN (60k params)
Full Transformer (109k params)
SCALE-Former (100k params)
Lower-left is preferable; latency is mean ± SD across five per-round medians
Held-out missing-region NMSE (dB)
−26
−27
−28
−29
−30
End-to-end latency (ms)
2
4
6
8
10
12
14